

# Actors

Paul Preißner

Fakultät für Informatik  
Technische Universität München

July 6<sup>th</sup> 2017

# Brief history


- ▶ C. Hewitt et al. '73 onward: first theory of actor model, operational semantics, axioms
- ▶ W. Clinger '81: proved unbounded nondeterminism property
- ▶ G. Agha '85: formalization of semantic model
- ▶ Theoretical/Practical research by MIT, CalTech, industry, etc.
- ▶ Recent resurgence (strong relevance to distributed/cloud computing)

# “A Model of Concurrent Computation in Distributed Systems”

- ▶ actors encapsulate computation (technically at any level)
- ▶ an actor may only send messages to actors it knows by name
- ▶ an (idling) actor receiving a message will accept it and execute the computation defined within, resulting in the possible actions:
  - ▶ sending new messages
  - ▶ creating new actors
  - ▶ updating its local state
- ▶ an actor can only influence its own local state

→ “self-contained, autonomous, interactive, asynchronously operating components” [Karmani, Agha]

## Example structure



**Figure:** Any actor may send messages to other known actors, create new actors or update its own state. [inspired by Karmani, Agha]

## (Example) Hello ...

using C++ Actor Framework; to illustrate:

```
1 [includes, usings]
2
3 behavior pong(event_based_actor* self, string selfname) {
4     return {
5         //if the message contains a string, proceed
6         [=](const string& what) -> string {
7             aout(self) << selfname << ":␣" << what << endl;
8             // reply Pong
9             return string("Pong!");
10        }
11    };
12 }
```

Specify behavior

## (Example) Hello ...

```
1 void ping(event_based_actor* self, const actor& buddy,  
2     string selfname) {  
3     // send Ping to buddy (timeout for reply = 10s)  
4     self->request(buddy, std::chrono::seconds(10),  
5         "Ping!").then(  
6         //if the message contains a string, proceed  
7         [=](const string& what) {  
8             aout(self) << selfname << ":_␣" << what << endl;  
9             //if reply is as expected, restart ping again  
10            if(what.compare("Pong...") == 0)  
11                ping(self, buddy, selfname);  
12        }  
13    );  
14 }
```

Specify actions

## (Example) ... World!

```
1 int main() {  
2     [caf setup]  
3     // create a new actor that calls 'pong()'  
4     auto actor_B = system.spawn(pong, "B");  
5     // create another actor that calls 'ping(actor_B)';  
6     auto actor_A = system.spawn(ping, actor_B, "A"); }
```

### Spawn actors and start something

```
1 B: Ping...  
2 A: Pong...  
3 B: Ping...  
4 A: Pong...  
5 B: Ping...  
6 A: Pong...  
7 [...]
```

Output (CAF "automatically" schedules to achieve high utilization across all 4 threads of i7 7500U)

# Main semantic properties

- ▶ Encapsulation & atomic execution:  
actors don't share state;  
process one message at a time, arrivals mid-computation need to be buffered;
- ▶ Fairness  
every actor makes progress, every message is delivered eventually;  
assumes fair scheduler;  
→actor cannot stall entire program
- ▶ Location transparency  
physical location not bound to identifier;  
→(hidden) migration possible, i.e. *mobility* →allows load-balancing, efficiency optimization



# Synchronization ...

Synchronization has to be through messages

- ▶ intuitively: "Remote Procedure Call"-like messaging  
send request, wait (and defer) until matching reply → allows  
e.g. ordering
- ▶ local synchronization constraints: specify constraints on actors  
by which they accept or deny messages
- ▶ synchronizers: similar to pre-processing of messages

→ use methods to construct complex work patterns

## ... and abstraction

- ▶ use patterns to structure actors into different compositions
- ▶ use abstract compositions to group actors and allow more efficient identification or message recipients

# Worst practices

- ▶ Faithful implementation of the model tends to be inefficient, shortcuts tend to be taken (that violate its properties)  
→ correctness of execution after optimization has to be checked
- ▶ Message latency: more distance  $\bar{m}$ ore latency  
→ use communication-computation overlap and suitable decomposition/migration to mask
- ▶ Naive send vs channels: individual sends may have high overhead  
→ utilize channels with stateful channel contracts to reduce overhead

# Worst practices

- ▶ Thread/switch overhead: the closer to a 1:1 map-to-thread mapping, the higher the overhead for switching execution to other actors; additional overhead for thread creation  
→ Use continuations based actors that don't perform full context switches and have reduced creation overhead
- ▶ Copying vs referencing: model demands no state be shared, naively this means no references be sent, only deep copies  
→ carefully allow send-by-reference for immutable types
- ▶ Scheduling: scheduler needs to guarantee fairness as in the model, not all schedulers satisfy this  
→ modify on a case-by-case basis; lazy thread creation - when tweaked - can be a relatively simple fix

# Support

- ▶ ~25 native actor languages
- ▶ >50 libraries for common languages (C, C++, C#, Java, JS, Python, Ruby, Haskell, LabVIEW, .NET, etc)
- ▶ most still actively supported, however some lacking proper documentation and developer support
- ▶ plug-ins for some IDEs exist (e.g. Erlang, Scala for Eclipse), some testing tools as well
- ▶ widespread support and documentation still spotty

# Distributed systems

DS server HPC microservices

# Embedded systems

system level appl design

# Versus

???



# Q&A