

Actors

Paul Preißner

Fakultät für Informatik
Technische Universität München

July 6th 2017

Brief history

- ▶ C. Hewitt et al. '73 onward: first theory of actor model, operational semantics, axioms
- ▶ W. Clinger '81: proved unbounded nondeterminism property
- ▶ G. Agha '85: formalization of semantic model
- ▶ Theoretical/Practical research by MIT, CalTech, industry, etc.
- ▶ Recent resurgence (strong relevance to distributed/cloud computing)

“A Model of Concurrent Computation in Distributed Systems”

- ▶ actors encapsulate computation (technically at any level)
- ▶ an actor may only send messages to actors it knows by name
- ▶ an (idling) actor receiving a message will accept it and execute the computation defined within, resulting in the possible actions:
 - ▶ sending new messages
 - ▶ creating new actors
 - ▶ updating its local state
- ▶ an actor can only influence its own local state

→ “self-contained, autonomous, interactive, asynchronously operating components” [Karmani, Agha]

Example structure




Figure: Any actor may send messages to other known actors, create new actors or update its own state. [inspired by Karmani, Agha]

(Example) Hello ...

using C++ Actor Framework; to illustrate:

```
1 [includes, usings]
2
3 behavior pong(event_based_actor* self, string selfname) {
4     return {
5         //if the message contains a string, proceed
6         [=](const string& what) -> string {
7             aout(self) << selfname << ":␣" << what << endl;
8             // reply Pong
9             return string("Pong!");
10        }
11    };
12 }
```

Specify behavior

(Example) Hello ...

```
1 void ping(event_based_actor* self, const actor& buddy,  
2     string selfname) {  
3     // send Ping to buddy (timeout for reply = 10s)  
4     self->request(buddy, std::chrono::seconds(10),  
5         "Ping!").then(  
6         //if the message contains a string, proceed  
7         [=](const string& what) {  
8             aout(self) << selfname << ":_␣" << what << endl;  
9             //if reply is as expected, restart ping again  
10            if(what.compare("Pong...") == 0)  
11                ping(self, buddy, selfname);  
12        }  
13    );  
14 }
```

Specify actions

(Example) ... World!

```
1 int main() {  
2     [caf setup]  
3     // create a new actor that calls 'pong()'  
4     auto actor_B = system.spawn(pong, "B");  
5     // create another actor that calls 'ping(actor_B)';  
6     auto actor_A = system.spawn(ping, actor_B, "A"); }
```

Spawn actors and start something

```
1 B: Ping...  
2 A: Pong...  
3 B: Ping...  
4 A: Pong...  
5 B: Ping...  
6 A: Pong...  
7 [...]
```

Output (CAF "automatically" schedules to achieve high utilization across all 4 threads of i7 7500U)

Main semantic properties

- ▶ Encapsulation & atomic execution:
 - actors don't share state;
 - process one message at a time, arrivals mid-computation need to be buffered;
- ▶ Fairness
 - every actor makes progress, every message is delivered eventually;
 - assumes fair scheduler;
 - actor cannot stall entire program
- ▶ Location transparency
 - physical location not bound to identifier;
 - (hidden) migration possible, i.e. *mobility* →allows load-balancing, efficiency optimization

Synchronization ...

Synchronization has to be through messages

- ▶ intuitively: "Remote Procedure Call"-like messaging
send request, wait (and defer) until matching reply →allows
e.g. ordering
- ▶ local synchronization constraints: specify constraints on actors
by which they accept or deny messages
- ▶ synchronizers: similar to pre-processing of messages

→use methods to construct complex work patterns

... and abstraction

- ▶ use patterns to structure actors into different compositions
- ▶ use abstract compositions to group actors and allow more efficient identification or message recipients

Worst practices

- ▶ Faithful implementation of the model tends to be inefficient, shortcuts tend to be taken (that violate its properties)
→ correctness of execution after optimization has to be checked
- ▶ Message latency: more distance = more latency
→ use communication-computation overlap and suitable decomposition/migration to mask
- ▶ Naive send vs channels: individual sends may have high overhead
→ utilize channels with stateful channel contracts to reduce overhead


Worst practices

- ▶ Thread/switch overhead: the closer to a 1:1 map-to-thread mapping, the higher the overhead for switching execution to other actors; additional overhead for thread creation
→ Use continuations based actors that don't perform full context switches and have reduced creation overhead
- ▶ Copying vs referencing: model demands no state be shared, naively this means no references be sent, only deep copies
→ carefully allow send-by-reference for immutable types
- ▶ Scheduling: scheduler needs to guarantee fairness as in the model, not all schedulers satisfy this
→ modify on a case-by-case basis; lazy thread creation - when tweaked - can be a relatively simple fix

Support

- ▶ ~25 native actor languages
- ▶ >50 libraries for common languages (C, C++, C#, Java, JS, Python, Ruby, Haskell, LabVIEW, .NET, etc)
- ▶ most still actively supported, however some lacking proper documentation and developer support
- ▶ plug-ins for some IDEs exist (e.g. Erlang, Scala for Eclipse), some testing tools as well
- ▶ widespread support and documentation still spotty

in Distributed systems



- ▶ Naive: map actors directly to clients
- ▶ More suitable: map actors to program parts, e.g. individual services, users, etc → e.g. *Facebook* chat, *Twitter* queues, *Halo 4* services, *Lift* web-app framework
- ▶ HPC: actor model highly scalable (with care), e.g. ActorX10 project at TUM
- ▶ Microservices: composition of independent services

in Embedded systems

- ▶ in system level design: model components of the system as actors → model-based design (capture requirements of ES)
→ enables better abstraction, definition of interface, reasoning over execution
- ▶ on application level: e.g. ActorX10 project, model stages of algorithm as actors, pipeline execution to achieve high utilization → optimization of stages possible

Versus

- ▶ Petri nets: places (conditions), transitions (events), arcs between places \leftrightarrow transitions \rightarrow net of pre-/postconditions of possible events
 - \rightarrow any number of tokens in the net, tokens trigger transitions and are transferred to postconditions
 - \rightarrow actors modeled as units of computation, petri nets as transitions/events mostly independent of units

Versus

- ▶ process calculi: collection of formal models, processes interact by message-passing (anonymous, through channels), processes as composition of primitives and operators subject to algebraic laws
 - message-passing similar to actors, processes could be modeled to resemble actors
- ▶ I/O automata: automaton modeled into any kind of single component, as state machine with in-/output and (hidden) internal actions
 - similar to encapsulation in actors, automata can be modeled to resemble actor model structure

Conclusion?

- ▶ actor model has great potential in achieving scalable concurrency
- ▶ quantitative and qualitative developer support is lacking, but increasing
- ▶ more and more use in large and small, distributed and embedded, commercial and free/open projects

Q&A